


```
• batches = Flux.Data.DataLoader(binarized_MNIST, batchsize=BS)
```

```
(784, 200)
```

```
• # confirm dimensions are as expected (D,BS)
• size(first(batches))
```

Bernoulli Log Density

The Bernoulli distribution $\text{Ber}(x \mid \mu)$ where $\mu \in [0, 1]$ is difficult to optimize for a few reasons. One solution is to parameterize the "logit-means": $y = \log(\frac{\mu}{1-\mu})$.

We can exploit further numerical stability, e.g. in computing $\log(1 + \exp(x))$, using library provided functions `log1pexp`

```
• using StatsFuns: log1pexp #log(1 + exp(x))
```

`bernoulli_log_density` (generic function with 1 method)

```
• # Numerically stable bernoulli density, why do we do this?
• function bernoulli_log_density(x, logit_means)
•     """Numerically stable log_likelihood under bernoulli by accepting  $\mu/(1-\mu)$ """
•     b = x .* 2 .- 1 # [0,1] -> [-1,1]
•     return - log1pexp.(-b .* logit_means)
• end
```

Model Implementation

- `log_prior` that computes the log-density of a latent representation under the prior distribution.
- `decoder` that takes a latent representation z and produces a 784-dimensional vector y . This will be a simple neural network with the following architecture: a fully connected layer with 500 hidden units and `tanh` non-linearity, a fully connected output layer with 784-dimensions. The output will be unconstrained, no activation function.
- `log_likelihood` that given an array binary pixels x and the output from the decoder, y corresponding to "logit-means" of the pixel Bernoullis $y = \log(\frac{\mu}{1-\mu})$ compute the **log**-likelihood under our model.
- `joint_log_density` that uses the `log_prior` and `log_likelihood` and gives the log-density of their joint distribution under our model $\log p_{\theta}(x, z)$.

Note that these functions should accept a batch of digits and representations, an array with elements concatenated along the last dimension.

`factorized_gaussian_log_density` (generic function with 1 method)

```
• function factorized_gaussian_log_density(samples,  $\mu$ ,  $\log\sigma$ )
•      $\sigma = \exp(\log\sigma)$ 
•     return sum(-0.5*((samples.- $\mu$ )./ $\sigma$ ).^2 .- log( $\sigma$ .*sqrt(2 $\pi$ )),dims=1)
• end
```

`log_prior` (generic function with 1 method)

```
• log_prior(z) = factorized_gaussian_log_density(z, 0, 0)
```

```
(2, 500, 784)
```

```
• Dz, Dh, Ddata = 2, 500, 28^2
```

```
decoder = Chain(Dense(2, 500, tanh), Dense(500, 784))
```

- `decoder = Chain(Dense(Dz, Dh, tanh), Dense(Dh, Ddata))` *# You can use Flux's Chain and Dense here*

```
log_likelihood (generic function with 1 method)
```

- `function log_likelihood(x,z)`
- `""" Compute log likelihood log_p(x|z) """`
- `# use numerically stable bernoulli`
- `return sum(bernoulli_log_density(x, decoder(z)), dims=1)`
- `end`

```
joint_log_density (generic function with 1 method)
```

- `joint_log_density(x,z) = log_prior(z) .+ log_likelihood(x,z)`

Amortized Approximate Inference with Learned Variational Distribution

Now that we have set up a model, we would like to learn the model parameters θ . Notice that the only indication for *how* our model should represent digits in $z \in \mathbb{R}^2$ is that they should look like our prior $\mathcal{N}(0, 1)$.

How should our model learn to represent digits by 2D latent codes? We want to maximize the likelihood of the data under our model $p_\theta(x) = \int p_\theta(x, z) dz = \int p_\theta(x | z) p(z) dz$.

We have learned a few techniques to approximate these integrals, such as sampling via MCMC. Also, 2D is a low enough latent dimension, we could numerically integrate, e.g. with a quadrature.

Instead, we will use variational inference and find an approximation $q_\phi(z) \approx p_\theta(z | x)$. This approximation will allow us to efficiently estimate our objective, the data likelihood under our model. Further, we will be able to use this estimate to update our model parameters via gradient optimization.

Following the motivating paper, we will define our variational distribution as q_ϕ also using a neural network. The variational parameters, ϕ are the weights and biases of this "encoder" network.

This encoder network q_ϕ will take an element of the data x and give a variational distribution over latent representations. In our case we will assume this output variational distribution is a fully-factorized Gaussian. So our network should output the $(\mu, \log \sigma)$.

To train our model parameters θ we will need also train variational parameters ϕ . We can do both of these optimization tasks at once, propagating gradients of the loss to update both sets of parameters.

The loss, in this case, no longer being the data likelihood, but the Evidence Lower BOund (ELBO).

1. Implement `log_q` that accepts a representation z and parameters $\mu, \log \sigma$ and computes the logdensity under our variational family of fully factorized gaussians.
2. Implement `encoder` that accepts input in data domain x and outputs parameters to a fully-factorized gaussian $\mu, \log \sigma$. This will be a neural network with fully-connected architecture, a

single hidden layer with 500 units and \tanh nonlinearity and fully-connected output layer to the parameter space.

3. Implement `elbo` which computes an unbiased estimate of the Evidence Lower BOund (using simple monte carlo and the variational distribution). This function should take the model p_θ , the variational model q_ϕ , and a batch of inputs x and return a single scalar averaging the ELBO estimates over the entire batch.
4. Implement simple loss function `loss` that we can use to optimize the parameters θ and ϕ with gradient . We want to maximize the lower bound, with gradient descent. (This is already implemented)

`log_q` (generic function with 1 method)

- `log_q(z, q_μ, q_logσ) = factorized_gaussian_log_density(z, q_μ, q_logσ)`

`unpack_guassian_params` (generic function with 1 method)

- `function unpack_guassian_params(output)`
- `μ, logσ = output[1:2,:], output[3:4,:]`
- `return μ, logσ`
- `end`

`encoder = Chain(Dense(784, 500, tanh), Dense(500, 4), unpack_guassian_params)`

- `encoder = Chain(Dense(Ddata, Dh, tanh), Dense(Dh, Dz*2), unpack_guassian_params)`

`sample_from_var_dist` (generic function with 1 method)

- `sample_from_var_dist(μ, logσ) = (randn(size(μ)) .* exp.(logσ) .+ μ)`

`elbo` (generic function with 1 method)

- `function elbo(x)`
- `#TODO variational parameters from data`
- `q_μ, q_logσ = encoder(x)`
- `#TODO: sample from variational distribution`
- `z = sample_from_var_dist(q_μ, q_logσ)`
- `#TODO: joint likelihood of z and x under model`
- `joint_ll = joint_log_density(x,z)`
- `#TODO: likelihood of z under variational distribution`
- `log_q_z = log_q(z, q_μ, q_logσ)`
- `#TODO: Scalar value, mean variational evidence lower bound over batch`
- `elbo_estimate = sum(joint_ll - log_q_z)/size(x)[2]`
- `return elbo_estimate`
- `end`

`loss` (generic function with 1 method)

- `function loss(x)`
- `return -elbo(x)`
- `end`

Optimize the model and amortized variational parameters

If the above are implemented correctly, stable numerically, and differentiable automatically then we can train both the `encoder` and `decoder` networks with graident optimization.

We can compute gradient `s` of our `loss` with respect to the `encoder` and `decoder` parameters `theta` and `phi` .

We can use a `Flux.Optimise` provided optimizer such as ADAM or our own implementation of gradient descent to update! the model and variational parameters.

Use the training data to learn the model and variational networks.

`train!` (generic function with 1 method)

```
• function train!(enc, dec, data; nepochs=100)
•   params = Flux.params(enc, dec)
•   opt = ADAM()
•   @info "Begin training in 2D latent space"
•   for epoch in 1:nepochs
•     b_loss = 0
•     for batch in data
•       # compute gradient wrt loss
•       grads = Flux.gradient(params) do
•         b_loss = loss(batch)
•       return b_loss
•     end
•     # update parameters
•     Flux.Optimise.update!(opt, params, grads)
•   end
•   # Optional: log loss using @info "Epoch $epoch: loss:..."
•   @info "Epoch $epoch: loss:$b_loss"
•   # Optional: visualize training progress with plot of loss
• end
• @info "Training in 2D is done"
• # return nothing, this mutates the parameters of enc and dec!
• end
```

```
• train!(encoder, decoder, batches, nepochs=5)
```

```
• using BSON: @save
```

```
• begin
•   @save "encoder.bson" encoder
•   @save "decoder.bson" decoder
• end
```

```
• using BSON
```

```
Dict(:decoder => Chain(Dense(2, 500, tanh), Dense(500, 784)))
```

```
• begin
•   BSON.load("encoder.bson", @__MODULE__)
•   BSON.load("decoder.bson", @__MODULE__)
• end
```

```
(2×200 Matrix{Float32}:
 -0.858582  0.00238673  1.66498 -4.57047 ... -1.82859  0.609256  0.0182087 -2.076
  0.608285  1.3631    -1.28454  0.668784 ...  0.646646  0.100764  0.646799 -2.117
```

```
• q_μ, q_logσ = encoder(first(batches))
```

Visualizing the Model Learned Representation

We will use the model and variational networks to visualize the latent representations of our data learned by the model.

We will use a variety of qualitative techniques to get a sense for our model by generating distributions over our data, sampling from them, and interpolating in the latent space.

- using **Plots** **##**

- using **Images**

calculate_bernoulli_mean (generic function with 1 method)

- function calculate_bernoulli_mean(logit_means)
- return exp.(logit_means) ./ (1 .+ exp.(logit_means))
- end

Larger Latent Space

Experimented a 3D latent space and make visualization

create_enc_dec (generic function with 1 method)

- function create_enc_dec(Dz, unpack_method)
- encoder = Chain(Dense(Ddata, Dh, tanh), Dense(Dh, Dz*2), unpack_method)
- decoder = Chain(Dense(Dz, Dh, tanh), Dense(Dh, Ddata))
- return encoder, decoder
- end

Dz_3d = 3

- Dz_3d = 3

unpack_guassian_params_3d (generic function with 1 method)

- function unpack_guassian_params_3d(output)
- μ , $\log\sigma$ = output[1:3,:], output[4:6,:]
- return μ , $\log\sigma$
- end

(Chain(Dense(784, 500, tanh), Dense(500, 6), unpack_guassian_params_3d), Chain(Dense(3,

- encoder_3d, decoder_3d = create_enc_dec(Dz_3d, unpack_guassian_params_3d)

log_likelihood_larger (generic function with 1 method)

- function log_likelihood_larger(x,z)
- """ Compute log likelihood log_p(x|z)"""
- return sum(bernoulli_log_density(x, decoder_3d(z)),dims=1)
- end

sample_from_var_dist_3d (generic function with 1 method)

- sample_from_var_dist_3d(μ , $\log\sigma$) = (randn(size(μ)) .* exp($\log\sigma$) .+ μ)

joint_log_density_3d (generic function with 1 method)

- joint_log_density_3d(x,z) = log_prior(z) .+ log_likelihood_larger(x,z)

elbo_3d (generic function with 1 method)

- function elbo_3d(x)
- q_ μ , q_ $\log\sigma$ = encoder_3d(x)
- z = sample_from_var_dist_3d(q_ μ , q_ $\log\sigma$)
- joint_ll = joint_log_density_3d(x,z)
- log_q_z = log_q(z, q_ μ , q_ $\log\sigma$)
- elbo_estimate = sum(joint_ll - log_q_z)/size(x)[2]
- return elbo_estimate, q_ $\log\sigma$
- end

loss_3d (generic function with 1 method)

- function loss_3d(x)
- elbo_estimate, $\log\sigma$ = elbo_3d(x)
- return -elbo_estimate, $\log\sigma$
- end

- *# logσ = Any[]*

train_3d! (generic function with 1 method)

```

• function train_3d!(enc, dec, data; nepochs=100)
•     params = Flux.params(enc, dec)
•     opt = ADAM()
•     @info "Begin training in 3D latent space"
•     for epoch in 1:nepochs
•         b_loss = 0
•         logσ = Any[]
•         for batch in data
•             grads = Flux.gradient(params) do
•                 b_loss, logσ = loss_3d(batch)
•                 # push!(logσ, logσ)
•                 # if epoch == nepochs
•                 #     var = (exp.(logσ)) .^ 2
•                 #     vs = size(var)
•                 # end
•                 return b_loss
•             end
•             Flux.Optimise.update!(opt, params, grads)
•         end
•         var = (exp.(logσ)) .^ 2
•         # var_size = size(var)
•         # @info "vs: $var_size"
•         scatter(var[1,:], var[2,:], var[3,:])
•         # Optional: log loss using @info "Epoch $epoch: loss:..."
•         @info "Epoch $epoch: loss:$b_loss"
•         # Optional: visualize training progress with plot of loss
•         # Optional: save trained parameters to avoid retraining later
•     end
•     @info "Training in 3D is done"
•     # return nothing, this mutates the parameters of enc and dec!
•     # return logσ
• end

```

- train_3d!(encoder_3d, decoder_3d, batches, nepochs=5)

```

(3×200 Matrix{Float32}:
  -0.0739408  0.943786  0.187911  -2.48795  ...  -0.646299  -1.39841  -0.0249961  -2.2
  -0.990187  -0.55024  3.25409  -3.99575  ...  -1.94905   0.237793  -1.2267    -2.1
   0.442513  -0.398217  0.211053  -0.910293  ...  -1.36731   3.3748    0.824378  -1.8
, 3×200

```

- q_μ_3d, q_logσ_3d = encoder_3d(first(batches))

```

[5, 0, 4, 1, 9, 2, 1, 3, 1, 4, 3, 5, 3, 6, 1, 7, 2, 8, 6, 9,      more , 2, 5


```

```

• # Training set labels
• begin
•     train_labels = Flux.Data.MNIST.labels(:train)
•     label_batches = Flux.Data.DataLoader(train_labels, batchsize=BS)
•     labels = first(label_batches)
• end

```

A batch 3D latent space of mean vectors for μ



```
• scatter(q_μ_3d[1,:), q_μ_3d[2,:), q_μ_3d[3,:), group=labels, title="A batch 3D latent
space of mean vectors for μ", xlabel="z1 for mean μ", ylabel="z2 for mean μ",
zlabel="z3 for mean μ")
```

Comparison with baselines

draw_image (generic function with 1 method)

```
• # Helper function for drawing the MNIST digit in 28*28 shape
• function draw_image(x)
•   dim = ndims(x)
•   if dim == 2
•     x_2d = reshape(x, 28, 28, :)
•     return Gray.(x_2d)
•   else
•     x_3d = reshape(x, 28, 28)
•     return Gray.(x_3d)
•   end
• end
```

visualize_samples (generic function with 1 method)

```
• function visualize_samples(decoder, dim)
•   plots1 = Any[]
•   plots = Any[]
•   for i in 1:5
•     # 1. Sample five 2D/3D zs from the prior p(z)
•     z = randn(dim,)
•     # 2. decode each z to get logit-means
•     logit_means = decoder(z)
•     # 3. Transfer logit-means to Bernoulli means μ
•     bern_mean = calculate_bernoulli_mean(logit_means)[1:784]
•     push!(plots, draw_image(bern_mean))
•     # 5. Sample 1 example from Bernoulli
•     samples1 = rand(Float64, size(bern_mean)) .< bern_mean
•     push!(plots1, draw_image(samples1))
•   end
•   return plots1, plots
• end
```

plot_mnist_image (generic function with 1 method)

```
• function plot_mnist_image(plots, plots1)
•   # 6. Display all plots in a single 10 x 4 grid
•   p = plot(layout = (5,1), size=(500,800))
```



```

•   for i in 1:5
•       heatmap!(cat(plots[i], plots1[i], dims=2), subplot=i)
•   end
•   plot(p)
• end

```

Baseline (2D latent space)



```

• plots1_2D, plots_2D = visualize_samples(decoder, 2)

```






Model with larger dimension (3D latent space)








```

• plots1_3D, plots_3D = visualize_samples(decoder_3d, 3)

```



- `plot_mnist_image(plots_2D, plots_3D)`



```
• plot_mnist_image(plots1_2D, plots1_3D)
```

Variance respond as the dimensionality of latent space increases

2D latent space v.s. 3D latent space


```
q_var_3d =
3x200 Matrix{Float32}:
 0.0102591  0.00566477  0.00709673  0.0256235  ...  0.0125992  0.0299519  0.00961012
 0.0141764  0.0113781  0.0282756  0.0695169  ...  0.0261679  0.0356457  0.0261307
 0.026577  0.0195623  0.0351642  0.153108   ...  0.0449009  0.020621  0.0375741
```

```
• q_var_3d = (exp.(q_logσ_3d)).^2
```

```
q_var_2d =
2x200 Matrix{Float32}:
 0.015712  0.00800832  0.00578431  0.119965  ...  0.0194409  0.00770532  0.0132927
 0.0144806  0.00882775  0.0235562  0.0331087  ...  0.0193194  0.011105  0.0155383
```

```
• q_var_2d = (exp.(q_logσ)).^2
```

Variance along each dimension in 3D latent space




```
• begin
•   plot(q_var_3d[1,:], label="3D z1", lw = 1.5)
•   plot!(q_var_3d[2,:], label="3D z2", lw = 1.5)
•   plot!(q_var_3d[3,:], label="3D z3", lw = 1.5)
•   xlabel!("Number of samples")
•   ylabel!("Variance")
•   title!("Variance along each dimension in 3D latent space")
• end
```

• using Statistics

```
3×1 Matrix{Float32}:
0.02445558
0.023458395
0.041988656
```

```
• mean(q_var_3d, dims=2)
```

Variance along each dimension in 2D latent space



```
• begin
•   plot(q_var_2d[1,:], label="2D z1", lw = 1.5)
•   plot!(q_var_2d[2,:], label="2D z2", lw = 1.5)
•   xlabel!("Number of samples")
•   ylabel!("Variance")
•   title!("Variance along each dimension in 2D latent space")
• end
```

```
2x1 Matrix{Float32}:
 0.022270106
 0.02870878
```

- `mean(q_var_2d, dims=2)`

As the dimension of latent space increases, the variance along each dimension tends to increase.

Condition on MNIST Digit Supervision

Horizontally concatenate labels to data

```
Chain(Dense(13, 500, tanh), Dense(500, 794))
```

```

• begin
•     encoder_cond = Chain(Dense(Ddata+10, Dh, tanh), Dense(Dh, (Dz_3d+10)*2),
unpack_gaussian_params_3d)
•     decoder_cond = Chain(Dense(Dz_3d+10, Dh, tanh), Dense(Dh, Ddata+10))
• end

```

Change labels to one-hot encoding vectors

- using Flux: `onehotbatch`

```
onehot_labels =
10x60000 Flux.OneHotArray{10,2,Vector{UInt32}}:
 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0
0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 1 0 0 1 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 1 0 0 0 0 0 0
0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 1 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1
0 0 0 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0
```



```

.   params = Flux.params(enc, dec)
.   opt = ADAM()
.   @info "Begin training in 3D latent space with given labels"
.   for epoch in 1:nepochs
.       b_loss = 0
.       for batch in data
.           grads = Flux.gradient(params) do
.               b_loss = loss_3d_cond(batch)
.               return b_loss
.           end
.           Flux.Optimise.update!(opt, params, grads)
.       end
.       @info "Epoch $epoch: loss:$b_loss"
.   end
.   @info "Training in 3D latent space(labels) is done"
. end

```

```

. train_cond!(encoder_cond, decoder_cond, batches_cond, nepochs=3)

```

```

(3×200 Matrix{Float32}:
 0.396789 -0.416867 -1.52017  3.08969 ... 1.92532  0.421166 -0.323954 , 3×200 Ma
-0.178393 -0.540125  0.37284  0.805868 ... 0.686274 -0.925214 -0.442824 -1.8285
 0.732969  1.88375  -0.598648  0.300632 ... 1.39632  -2.22553  0.781096 -1.4836

```


```

. q_μ_cond, q_logσ_cond = encoder_cond(first(batches_cond))

```

Visualize latent representation

A batch 3D latent space of mean vectors for μ | labels



```

. scatter(q_μ_cond[1,:], q_μ_cond[2,:], q_μ_cond[3:], group=labels, title="A batch 3D
  latent space of mean vectors for μ | labels", xlabel="z1 for mean μ", ylabel="z2 for
  mean μ", zlabel="z3 for mean μ")

```


draw_image_cond (generic function with 1 method)

```

. # Helper function for drawing the MNIST digit in 28*28 shape
. function draw_image_cond(x)
.     dim = ndims(x)
.     if dim == 2
.         x_2d = reshape(x, 28, 28, :)
.         return Gray.(x_2d)
.     else
.         x_3d = reshape(x, 28, 28)
.         return Gray.(x_3d)
.     end

```

```
• plots1_cond, plots_cond = visualize_samples(decoder_cond, 13)
```



- `plot_mnist_image(plots_cond, plots1_cond)`

Semi-supervised learning

[illegible]


```

0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0

```

```

• begin
•   onehot_semi = Matrix(onehot_labels)
•   index = rand(1:60000,30000)
•   onehot_semi[:,index] = zeros(10, 30000)
• end

```

batches_semi =

```

DataLoader(794x60000 BitMatrix:
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  ⋮           ⋮           ⋮           ⋮           ⋮           ⋮
  0 0 0 0 0 0 0 0 0 0 1 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  1 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 1
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 1 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 1 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 1 0 0 0 0 0

```

```

• batches_semi = Flux.Data.DataLoader(cat(binarized_MNIST, onehot_semi, dims=1),
  batchsize=BS)

```

Chain(Dense(13, 500, tanh), Dense(500, 794))

```

• begin
•   encoder_semi = Chain(Dense(Ddata+10, Dh, tanh), Dense(Dh, (Dz_3d+10)*2),
  unpack_guassian_params_3d)
•   decoder_semi = Chain(Dense(Dz_3d+10, Dh, tanh), Dense(Dh, Ddata+10))
• end

```

log_likelihood_semi (generic function with 1 method)

```

• function log_likelihood_semi(x,z)
•   """ Compute log likelihood log_p(x|z,c)"""
•   # Let's just label all samples to be digit 0
•   digit_0_zeroes = zeros((9,200))
•   digit_0_ones = ones((1,200))
•   digit_0 = cat(digit_0_ones, digit_0_zeroes, dims=1)
•   cond_z = cat(z, digit_0, dims=1)
•   return sum(bernoulli_log_density(x, decoder_semi(cond_z)),dims=1)
• end

```

joint_log_density_semi (generic function with 1 method)

```

• joint_log_density_semi(x,z) = log_prior(z) .+ log_likelihood_semi(x,z)

```

elbo_3d_semi (generic function with 1 method)

```

• function elbo_3d_semi(x)
•   q_μ, q_logσ = encoder_semi(x)
•   z = sample_from_var_dist_3d(q_μ, q_logσ)
•   joint_ll = joint_log_density_semi(x,z)
•   log_q_z = log_q(z, q_μ, q_logσ)
•   elbo_estimate = sum(joint_ll - log_q_z)/size(x)[2]
•   return elbo_estimate
• end

```

loss_3d_semi (generic function with 1 method)

```

• function loss_3d_semi(x)
•   return -elbo_3d_semi(x)


```

- end






train_semi! (generic function with 1 method)

```
• function train_semi!(enc, dec, data; nepochs=100)
•   params = Flux.params(enc, dec)
•   opt = ADAM()
•   @info "Begin training in 3D latent space with given semi labels"
•   for epoch in 1:nepochs
•     b_loss = 0
•     for batch in data
•       grads = Flux.gradient(params) do
•         b_loss = loss_3d_semi(batch)
•         return b_loss
•       end
•       Flux.Optimise.update!(opt, params, grads)
•     end
•     @info "Epoch $epoch: loss:$b_loss"
•   end
•   @info "Training in 3D latent space(semi labels) is done"
• end
```

- train_semi!(encoder_semi, decoder_semi, batches_semi, nepochs=3)



- plots1_semi, plots_semi = visualize_samples(decoder_semi, 13)



```
• plot_mnist_image(plots_semi, plots1_semi)
```

Optimizing Different Divergences

```
(Chain(Dense(784, 500, tanh), Dense(500, 4), unpack_guassian_params), Chain(Dense(2, 500
```

```
• encoder_js, decoder_js = create_enc_dec(2, unpack_guassian_params)
```

```
log_likelihood_js (generic function with 1 method)
```

```
• function log_likelihood_js(x,z)
•     """ Compute log likelihood log_p(x|z) """
•     return sum(bernoulli_log_density(x, decoder_js(z)), dims=1)
• end
```

```
joint_log_density_js (generic function with 1 method)
```

```
• joint_log_density_js(x,z) = log_prior(z) .+ log_likelihood_js(x,z)
```

elbo_js (generic function with 1 method)

```
• function elbo_js(x)
•   q_μ, q_logσ = encoder_js(x)
•   z = sample_from_var_dist(q_μ, q_logσ)
•   joint_ll = joint_log_density_js(x,z)
•   log_q_z = log_q(z, q_μ, q_logσ)
•
•   p = exp.(joint_ll)
•   q = exp.(log_q_z)
•   m = 0.5*log.(p+q)
•
•   pm = sum(joint_ll - m)/size(x)[2]
•   qm = sum(log_q_z - m)/size(x)[2]
•
•   elbo_estimate = 0.5*pm + 0.5*qm
•   return elbo_estimate
• end
```

loss_js (generic function with 1 method)

```
• function loss_js(x)
•   return -elbo_js(x)
• end
```






train_js! (generic function with 1 method)

```
• function train_js!(enc, dec, data; nepochs=100)
•   params = Flux.params(enc, dec)
•   opt = ADAM()
•   @info "Begin training in 2D latent space using JS Divergence"
•   for epoch in 1:nepochs
•     b_loss = 0
•     for batch in data
•       # compute gradient wrt loss
•       grads = Flux.gradient(params) do
•         b_loss = loss_js(batch)
•       end
•       # update parameters
•       Flux.Optimise.update!(opt, params, grads)
•     end
•     @info "Epoch $epoch: loss:$b_loss"
•   end
•   @info "Training in 2D using JS Divergence is done"
• end
```

```
• train_js!(encoder_js, decoder_js, batches, nepochs=3)
```



```
• plots1_js, plots_js = visualize_samples(decoder_js, 2)
```



```
• plot_mnist_image(plots_js, plots1_js)
```

More Expressive Likelihood Model

Use beta likelihood model with $\alpha = 2$, $\beta = 2$ on float MNIST

```
• using Distributions
```

```
• using SpecialFunctions
```

```
float_MNIST =
784×60000 Matrix{Float64}:
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```

0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0

```

```
• float_MNIST = convert(Array{Float64}, greyscale_MNIST)
```

beta_log_density (generic function with 1 method)

```

• function beta_log_density(x, logit_means, α, β)
•     b = x .* 2 .- 1
•     B = gamma(α)*gamma(β) / gamma(α+β)
•     return - log1pexp.(-b .* logit_means / B)
• end

```

log_likelihood_beta (generic function with 1 method)

```

• function log_likelihood_beta(x, z, α, β)
•     """ Compute log likelihood log_p(x|z) """
•     return sum(beta_log_density(x, decoder_beta(z), α, β), dims=1)
• end

```

joint_log_density_beta (generic function with 1 method)

```
• joint_log_density_beta(x,z, α, β) = log_prior(z) .+ log_likelihood_beta(x,z,α, β)
```

```
(Chain(Dense(784, 500, tanh), Dense(500, 4), unpack_guassian_params), Chain(Dense(2, 500
```

```
• encoder_beta, decoder_beta = create_enc_dec(2, unpack_guassian_params)
```

elbo_beta (generic function with 1 method)

```

• function elbo_beta(x)
•     q_μ, q_logσ = encoder_beta(x)
•     z = sample_from_var_dist(q_μ, q_logσ)
•     joint_ll = joint_log_density_beta(x,z, 2, 2)
•     log_q_z = log_q(z, q_μ, q_logσ)
•     elbo_estimate = mean(joint_ll - log_q_z)
•     return elbo_estimate
• end

```

loss_beta (generic function with 1 method)

```

• function loss_beta(x)
•     return -elbo_beta(x)
• end

```

train_beta! (generic function with 1 method)

```

• function train_beta!(enc, dec, data; nepochs=100)
•     params = Flux.params(enc, dec)
•     opt = ADAM()
•     @info "Begin training in 2D latent space using Beta likelihood on float MNIST"
•     for epoch in 1:nepochs
•         b_loss = 0
•         for batch in data
•             # compute gradient wrt loss
•             grads = Flux.gradient(params) do
•                 b_loss = loss_beta(batch)
•             end
•             return b_loss
•         end
•         # update parameters
•         Flux.Optimise.update!(opt, params, grads)
•     end
•     @info "Epoch $epoch: loss:$b_loss"
• end
• @info "Training in 2D using Beta likelihood is done"
• end

```

```
float_batches =
```

```
DataLoader(784x60000 Matrix{Float64}):
```






```
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
• float_batches = Flux.Data.DataLoader(float_MNIST, batchsize=BS)
```

```
• train_beta!(encoder_beta, decoder_beta, float_batches, nepochs=5)
```



```
• plots1_beta, plots_beta = visualize_samples(decoder_beta, 2)
```



```
• plot_mnist_image(plots_beta, plots1_beta)
```

Inference

Use the baseline model to infer the bottom of a digit given the top

```
Dhalf = 392
```

```
• Dhalf = Int(28*28/2)
```

draw_top_half_image (generic function with 1 method)

```
• # Helper function for drawing only top half of the MNIST digit in 28*28 shape
• function draw_top_half_image(x)
•     x = reshape(x, 28, 28, :)
•     bot_x = x[1:14,:,:]
•     return reshape(bot_x, (14,28))
• end
```


draw_bot_half_image (generic function with 1 method)

```
• # Helper function for drawing only top half of the MNIST digit in 28*28 shape
• function draw_bot_half_image(x)
•     x = reshape(x, 28, 28, :)
•     bot_x = x[15:28,:,:]
•     return reshape(bot_x, (14,28))
• end
```

log_p_top_z (generic function with 1 method)

```
• # Calculate log likelihood log_p(top/z)
• function log_p_top_z(top, z)
•     x_half = decoder(z)[1:Dhalf,:]
•     return sum(bernoulli_log_density(top,x_half),dims=1)
• end
```

Log joint density $p(z, top)$

joint_log_density_top (generic function with 1 method)

```
• # Calculate log_p(top, z)
• joint_log_density_top(top,z) = log_prior(z) .+ log_p_top_z(top,z)
```

Stochastic variational inference $p(z|top)$

```
q_μ_top = 2×1 Matrix{Float64}:
  0.6707227541019384
 -0.12684732107979857
```

```
• q_μ_top = randn(Dz,1)
```

```
q_logσ_top = 2×1 Matrix{Float64}:
 -1.6022092621845037
  0.7220838252581145
```

```
• q_logσ_top = randn(Dz,1)
```

elbo_top (generic function with 1 method)

```
• function elbo_top(top, q_μ_top, q_logσ_top)
•     z = sample_from_var_dist(q_μ_top, q_logσ_top)
•     joint_ll = joint_log_density_top(top,z)
•     log_q_z = log_q(z, q_μ_top, q_logσ_top)
•     elbo_estimate = mean(joint_ll - log_q_z)
•     return -elbo_estimate
• end
```

loss_top (generic function with 1 method)

```
• function loss_top(top, q_μ_top, q_logσ_top)
•     return -elbo_top(top, q_μ_top, q_logσ_top)
• end
```

n = 60000

```
• n = size(train_labels)[1]
```

indices_0 =

[2, 22, 35, 38, 52, 57, 64, 69, 70, 76, 82, 89, 96, 109, 115, 119, 120, 122,

```
• # Construct a dataset consists of digit 0
• indices_0 = [i for i in 1:n if train_labels[i]==0]
```

digit_0 =

784×5923 BitMatrix:

```
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0    0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0    0 0 0 0 0 0 0 0 0 0 0 0 0
```

```
digit_0 = binarized_MNIST[:,indices_0]
```


```

• function train_top!(q_μ, q_logσ, data, loss_func; nepochs=100)
•     params = Flux.params(q_μ, q_logσ)
•     opt = ADAM()
•     @info "Begin training to optimize q_μ and q_logσ"
•     for epoch in 1:nepochs
•         b_loss = 0
•         grads = Flux.gradient(params) do
•             b_loss = loss_func(data[1:Dhalf])
•             return b_loss
•         end
•         Flux.Optimise.update!(opt, params, grads)
•         @info "Epoch $epoch: loss:$b_loss"
•     end
•     @info "Optimizing q_μ and q_logσ is done"
• end

```

- `loss_tophalf(top) = loss_top(top, q_μ_top, q_logσ_top)`



- `size(digit_0)`



```
• train_top!(q_μ_top, q_logσ_top, digit_0[:,2], loss_tophalf, nepochs=2)
```

Original digit 0

Inferred digit 0



```

• begin
•   p_top = plot(layout = (1,2))
•
•   # Take a sample z from the approximate posterior
•   z_top = sample_from_var_dist(q_μ_top, q_logσ_top)
•   # Feed z to decoder
•   logits_mean_top = decoder(z_top)
•   # Convert to bernoulli mean
•   bern_mean_top = calculate_bernoulli_mean(logits_mean_top)
•   bot_part = draw_bot_half_image(bern_mean_top)
•   top_part = draw_top_half_image(test_img)
•   cat_img = cat(top_part, bot_part, dims=1)
•
•   # Plot original and inferred results
•   plot!(test_img, title="Original digit 0", subplot=1)
•   plot!(draw_image(vec(cat_img)), title= "Inferred digit 0", subplot=2)
•
•   plot(p_top)
• end

```

More interesting data

Train the VAE model on Fashion MNIST dataset

`train_fashion =`



(a vector displayed as a row to save space)

```
• train_fashion = Flux.Data.FashionMNIST.images(:train)
```

`greyscale_fashion =`

```
• greyscale_fashion = hcat(float.(reshape.(train_fashion,:))...)
```

```
binarized_fashion =
```

```
784x60000 BitMatrix:
```

```
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
⋮           ⋮           ⋮           ⋮           ⋮           ⋮           ⋮
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
```

```
• binarized_fashion = greyscale_fashion .> 0.5
```

```
fashion_batches =
```

```
DataLoader(784x60000 BitMatrix:
```

```
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
⋮           ⋮           ⋮           ⋮           ⋮           ⋮           ⋮
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0 0 0 0
```

```
• fashion_batches = Flux.Data.DataLoader(binarized_fashion, batchsize=BS)
```

```
(Chain(Dense(784, 500, tanh), Dense(500, 6), unpack_guasssian_params_3d), Chain(Dense(3,
```

```
• encoder_fashion, decoder_fashion = create_enc_dec(Dz_3d, unpack_guasssian_params_3d)
```

```
log_likelihood_fashion (generic function with 1 method)
```

```
• function log_likelihood_fashion(x,z)
•   return sum(bernoulli_log_density(x, decoder_fashion(z)),dims=1)
• end
```

```
joint_log_density_fashion (generic function with 1 method)
```

```
• joint_log_density_fashion(x,z) = log_prior(z) .+ log_likelihood_fashion(x,z)
```

```
elbo_fashion (generic function with 1 method)
```

```
• function elbo_fashion(x)
•   q_μ, q_logσ = encoder_fashion(x)
•   z = sample_from_var_dist_3d(q_μ, q_logσ)
•   joint_ll = joint_log_density_fashion(x,z)
•   log_q_z = log_q(z, q_μ, q_logσ)
•   elbo_estimate = mean(joint_ll - log_q_z)
•   return elbo_estimate
• end
```

```
loss_fashion (generic function with 1 method)
```

```
• function loss_fashion(x)
•   return -elbo_fashion(x)
• end
```

```
train_fashion! (generic function with 1 method)
```

```
• function train_fashion!(enc, dec, data; nepochs=100)
```

```


•   params = Flux.params(enc, dec)
•   opt = ADAM()
•   @info "Begin training on FashionMNIST in 3D latent space"
•   for epoch in 1:nepochs
•       b_loss = 0
•       for batch in data
•           grads = Flux.gradient(params) do
•               b_loss = loss_fashion(batch)
•               return b_loss
•           end
•           Flux.Optimise.update!(opt, params, grads)
•       end
•       @info "Epoch $epoch: loss:$b_loss"
•   end
•   @info "Training on FashionMNIST in 3D latent space is done"
• end

```

```

• train_fashion!(encoder_fashion, decoder_fashion, fashion_batches, nepochs=5)






```



```

• plots1_fashion, plots_fashion = visualize_samples(decoder_fashion, 3)

```



```
• plot_mnist_image(plots_fashion, plots1_fashion)
```

